

# Project Large-scale Data Engineering (LDE) Kick-off Meeting

**Dr.-Ing. Patrick Damme**

Technische Universität Berlin

Berlin Institute for the Foundations of Learning and Data

Big Data Engineering (DAMS Lab)

- **Hybrid Setting with Optional Attendance**

- In-person in MAR 0.015
- Virtual via zoom

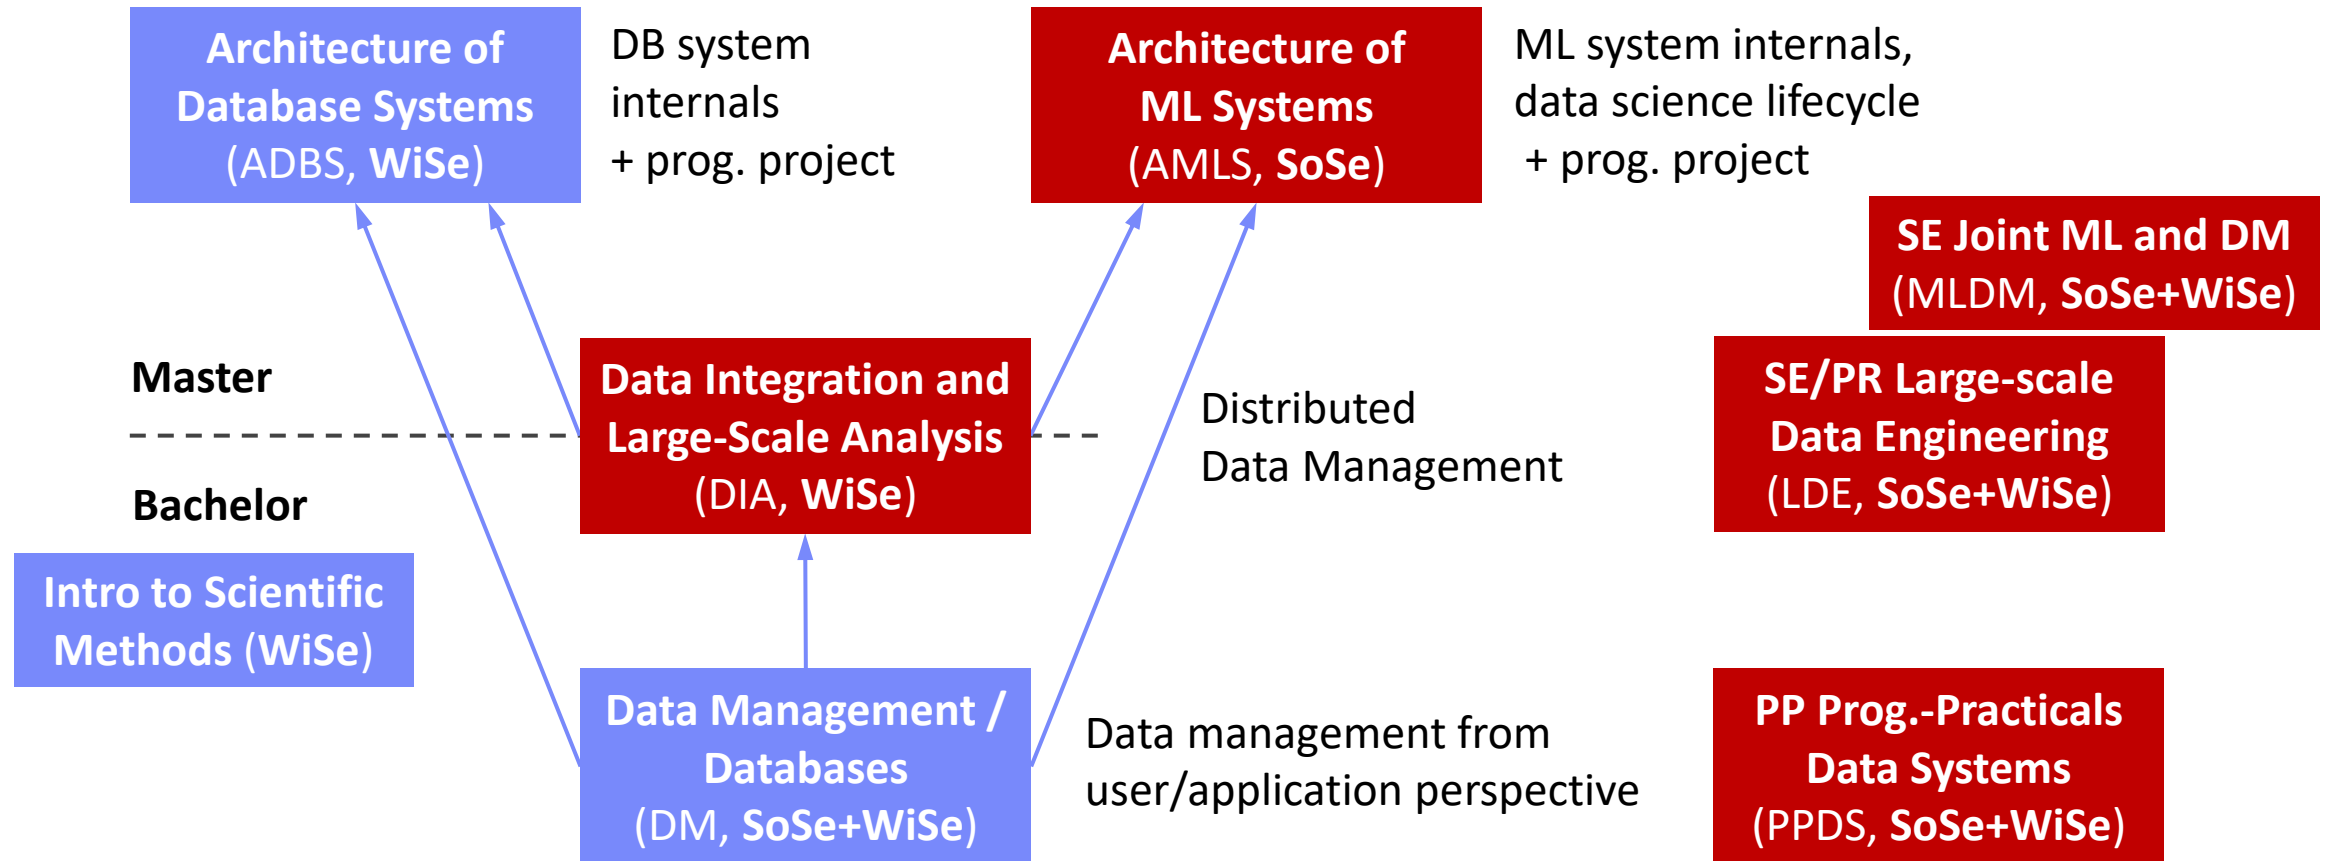


<https://tu-berlin.zoom-x.de/j/67376691490?pwd=NmlvWTM5VUVWRjU0UGI2bXhBVkxzQT09>

# About Me

- **Since 10/2022: Postdoc at TU Berlin, Germany**
  - FG Big Data Engineering (DAMS Lab) headed by Prof. Matthias Böhm
  - Continuing work on integrated data analysis pipelines
  - Research interests in the fields of database and ML systems (especially compiler & runtime techniques, extensibility)
- **2021-2022: Postdoc at TU Graz & Know-Center GmbH, Austria**
  - Data Management group headed by Prof. Matthias Böhm
  - Started work on integrated data analysis pipelines
- **2015-2020: PhD student at TU Dresden, Germany**
  - Dresden Database Research Group headed by Prof. Wolfgang Lehner
  - PhD thesis on making complex analytical database queries more efficient through lightweight compression of intermediate results





# Agenda



- **Course Organization, Outline, and Deliverables**
- **Projects in DAPHNE and Apache SystemDS**
- **How to Approach the Project**
- **List of Project Topics (Proposals)**

# Course Organization, Outline, and Deliverables

# Large-scale Data Engineering: Module Overview



20 seats in total

bachelor + master

#41086: LDE Seminar + Project (12 ECTS)

11 students

9 students

#41095: Seminar LDE (3)

#41183: Project LDE (9 ECTS)

8 students

bachelor-only

bachelor-only

Mon, 14:00-16:00  
MAR 0.015 & zoom

## Seminar LDE

- Reading & writing scientific papers
- Giving presentations on papers
- Summary paper
- Presentation
- Lecturer & seminar mentor



## Project LDE

- Building & evaluating prototypes
- Giving presentations on prototypes
- Prototype design/impl/tests/doc/eval
- Presentation
- Project mentors



Mon, 16:00-18:00  
MAR 0.015 & zoom

- In the context of systems for data engineering, data management, machine learning
- In combination: Ideal preparation for a bachelor/master thesis with our group

## ■ General Contact Person

- Dr.-Ing. Patrick Damme ([patrick.damme@tu-berlin.de](mailto:patrick.damme@tu-berlin.de))

## ■ Course Website

- [https://pdamme.github.io/teaching/2025-26\\_winter/lde/lde\\_winter2025-26.html](https://pdamme.github.io/teaching/2025-26_winter/lde/lde_winter2025-26.html)
- One site for seminar and project
- All material, schedule, **deadlines**

## ■ ISIS course

- <https://isis.tu-berlin.de/course/view.php?id=44129>
- Announcements, discussion forum, topic selection poll, submission of summary paper and presentation slides

## ■ Language

- Lectures and slides: **English**
- Communication: **English/German**
- Submitted paper and presentation: **English**
- **Informal language** (first name is fine), immediate feedback is welcome



# Semester Schedule & Deadlines



- **Kick-off Meeting Oct 13** (optional)
- **Recommended Introductory Lecture** (optional)
  - Oct 27, 14:00: Experiments, Reproducibility, and Giving Presentations
- **Self-organized Project Work**
  - Consultation hours for any questions (optional)
- **Intermediate Presentations** (prerequisite)
  - Jan 19, 16:00-18:00, MAR 0.009: All students
- **Final Presentations** (mandatory)
  - Feb 23, 14:00-18:00, MAR 0.009: All students
- **List of Project Topics**
  - Presented today, take your time to select afterwards
- **Topic Selection**
  - **Deadline: Oct 31, 23:59** (in 2½ weeks)
  - Ranked list of **5 topics** via poll on the ISIS course + pref on individual/team work [+ team members]
  - Global topic assignment based on preferences
  - **Notification of assigned topics: Nov 10** (in 4 weeks)
- **Submission of Initial Prototype** (prerequisite)
  - Implementation and tests
  - **Deadline: Jan 18, 23:59** (in 14 weeks)
  - As a pull request on GitHub (exceptionally by email)
- **Submission of Final Prototype** (mandatory)
  - Implementation, tests, docs, experiments
  - **Deadline: Feb 16, 23:59** (in 18 weeks)
  - As a pull request on GitHub (exceptionally by email)
- **Submission of Pres. Slides (Intermediate & Final Pres.)**
  - **Deadline: The day before the presentation, 23:59**
  - Upload PDF in the ISIS course

# Project Deliverables: Initial Prototype & Intermediate Presentation



- Introduced in **Response to Students' Feedback** (Course Evaluation)

- **Initial Prototype**

- 80% functionally complete prototype including good set of test cases
- Basis for further improvements driven by experiments and feedback

- **Intermediate Presentation**

- Slide presentation of 5-10 min per individual/team
- Briefly **present the problem** you work on
- Give an **overview of your initial prototype** (concepts and crucial changes to the code base)
- Outline your **planned experiments**
- Should be the result of **prior discussions with your project mentor**

**Ungraded Prerequisites  
for the Portfolio Exam**

to be allowed make mistakes  
and learn from them

- **Benefits for You**

- **Improved time management** (retain enough time for experiments)
- **Exchange with the other students** in the project
- **Get feedback** by project mentors and other students for improving the **quality of your prototype**

# Project Deliverables: Final Prototype & Final Presentation



## ■ Final Prototype

- 100% functionally complete prototype including good set of test cases
- Efficiency confirmed by experiments

## ■ Final Presentation

- Summarize the problem and give an overview of your final prototype
- Present your experimental results
- 1 student: 10 min talk + 5 min discussion = 15 min
- 2 students: 13 min + 7 min = 20 min
- 3 students: 16 min + 9 min = 25 min
- Audience: engage in the discussion

## ■ Grading

### ■ #41086 (seminar + project)

- Graded portfolio exam
- 25 pts: summary paper
- 15 pts: presentation
- 50 pts: design/impl/tests/doc
- 10 pts: presentation

### ■ #41183 (project-only)

- Graded portfolio exam
- 85 pts: implementation/tests/documentation
- 15 pts: presentation

## ■ Academic Honesty / No Plagiarism

implies that use of LLMs like ChatGPT is prohibited

# Portfolio Exam Registration



- **Portfolio exam registration: Nov 10 – Dec 08**
  - Binding registration in Moses/MTS
  - Including selection of seminar presentation date (first-come-first-serve)
- **Portfolio exam de-registration**
  - **Until 3 days before the first graded exam part**
    - Modules “LDE”/”Seminar LDE”: until **Jan 09**
    - Module “Project LDE”: until **Feb 13**
    - De-register yourself in Moses/MTS
  - **With sufficient reason: Until the day of the exam**
    - In case of sickness etc.
    - Modules “LDE”/”Seminar LDE”: until **Jan 12/Jan 26/Feb 02**
    - Module “Project LDE”: until **Feb 16/Feb 23**
- **Missing deadlines/exam without de-registration**
  - Zero points in the respective exam part (!)
  - **Approach us early in case of problems**
- **If you don't want to take LDE anymore**
  - Let me know asap to give students in the queue a chance to fill in

# LDE Project Goals and Mindset



## ■ Goals

- **Design/implement a prototype** in DAPHNE/SystemDS
- **AND prove it is a valuable contribution** to the system (tests, documentation, experiments)
- **Present and defend your work** in a presentation & discussion

**High-quality and  
convincing contribution  
to an open-source system**

## ■ Focus on Methodology

- LDE as a preparation for a bachelor/master thesis at the DAMS Lab

## ■ Mindset: Be Open Learn New Things and to Work on Any Part of the System

- Whatever it takes to fulfill the task
- Self-guided acquisition of required technical skills

## ■ Grading Criteria

- Design/implementation (functionality)
- Code quality + tests + documentation
- Experiments

# Projects in DAPHNE and Apache SystemDS

# Overview: Two Systems Developed at the DAMS Lab



## ■ An Open and Extensible System Infrastructure for **Integrated Data Analysis Pipelines**

- Intersection of data management, machine learning, high-performance computing
- Open-source (Apache v2 license)  
<https://github.com/daphne-eu/daphne>
- Originated from DAPHNE EU-project
- Written mostly in C++, Python, DaphneDSL
- Since 2020 (open-source since 2022)



[Patrick Damme et al.: DAPHNE: An Open and Extensible System Infrastructure for Integrated Data Analysis Pipelines. CIDR 2022]



## ■ A Declarative ML System for the **End-to-End Data Science Lifecycle**

- Data integration/cleaning/prep, model selection/training/validation/debugging/deployment/scoring
- Open-source (Apache v2 license)  
<https://github.com/apache/systemds>
- Originated from Apache SystemML (started at IBM)
- Written mostly in Java, Python, and DML
- Since 2010 (open-source since 2015) (as SystemML)



[Matthias Boehm et al.: SystemDS: A Declarative Machine Learning System for the End-to-End Data Science Lifecycle. CIDR 2020]

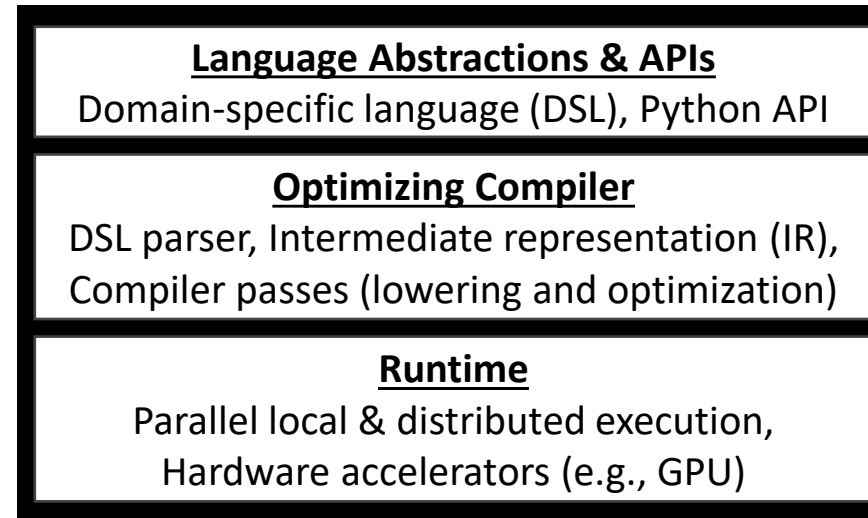

# Simplified High-level Architecture of DAPHNE and SystemDS



clustering ... neural networks  
classification ... regression ...  
relational queries

**Program**  
matrix/frame data types  
linear/relational algebra ops  
complex control flow

**Input Data**  
open file formats  
data exchange with other systems/libs



**Output Data**

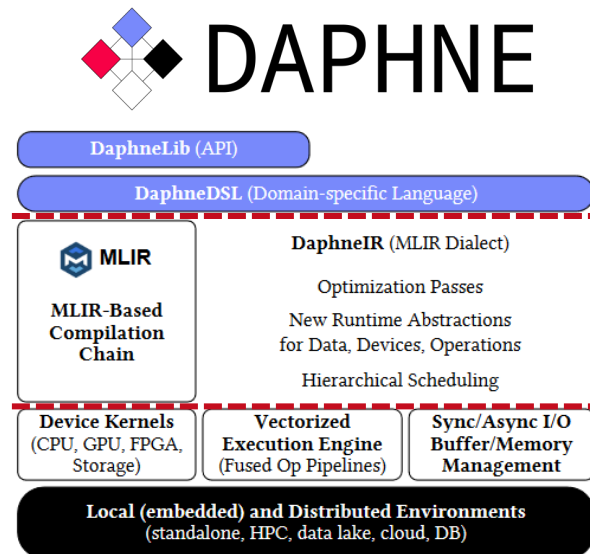
**Other such systems (examples)**



**Architecture of ML Systems  
(AMLS): Lecture + Exercise  
by Prof. Matthias Böhm  
(every summer semester)**

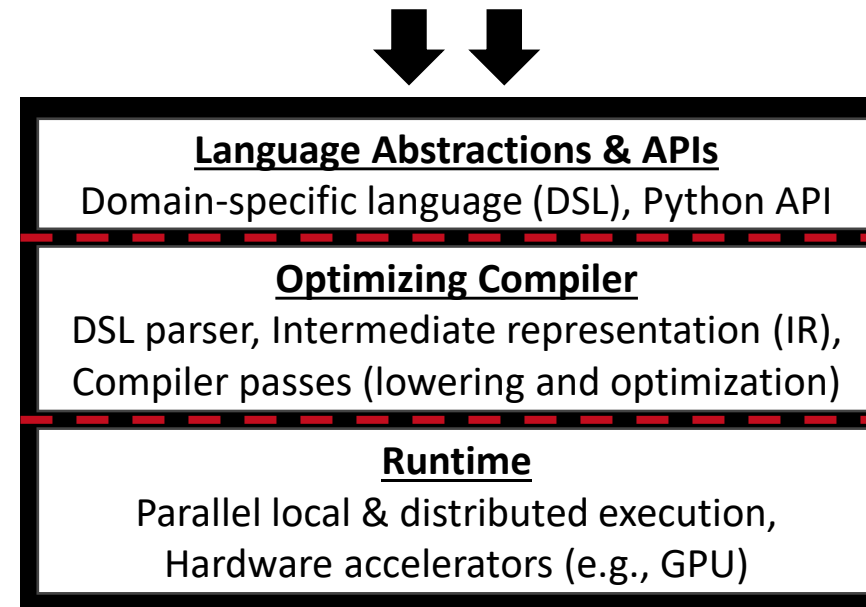


# Simplified High-level Architecture of DAPHNE and SystemDS

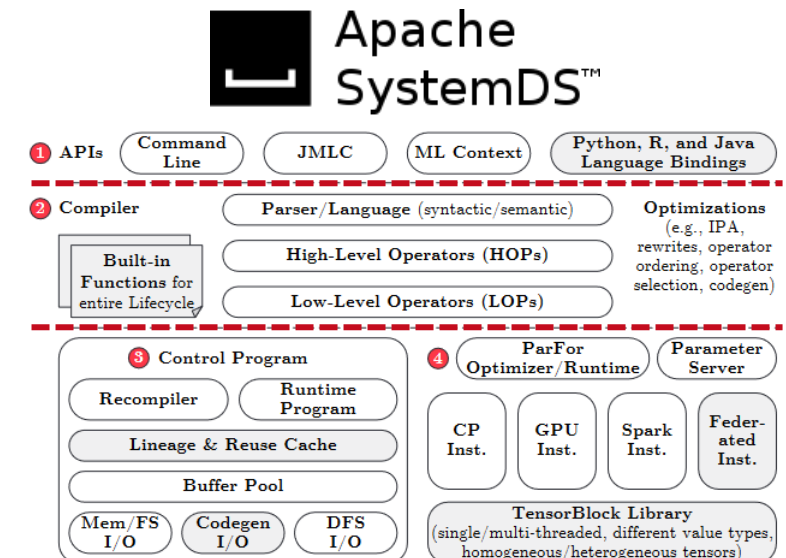


**Program**  
matrix/frame data types  
linear/relational algebra ops  
complex control flow

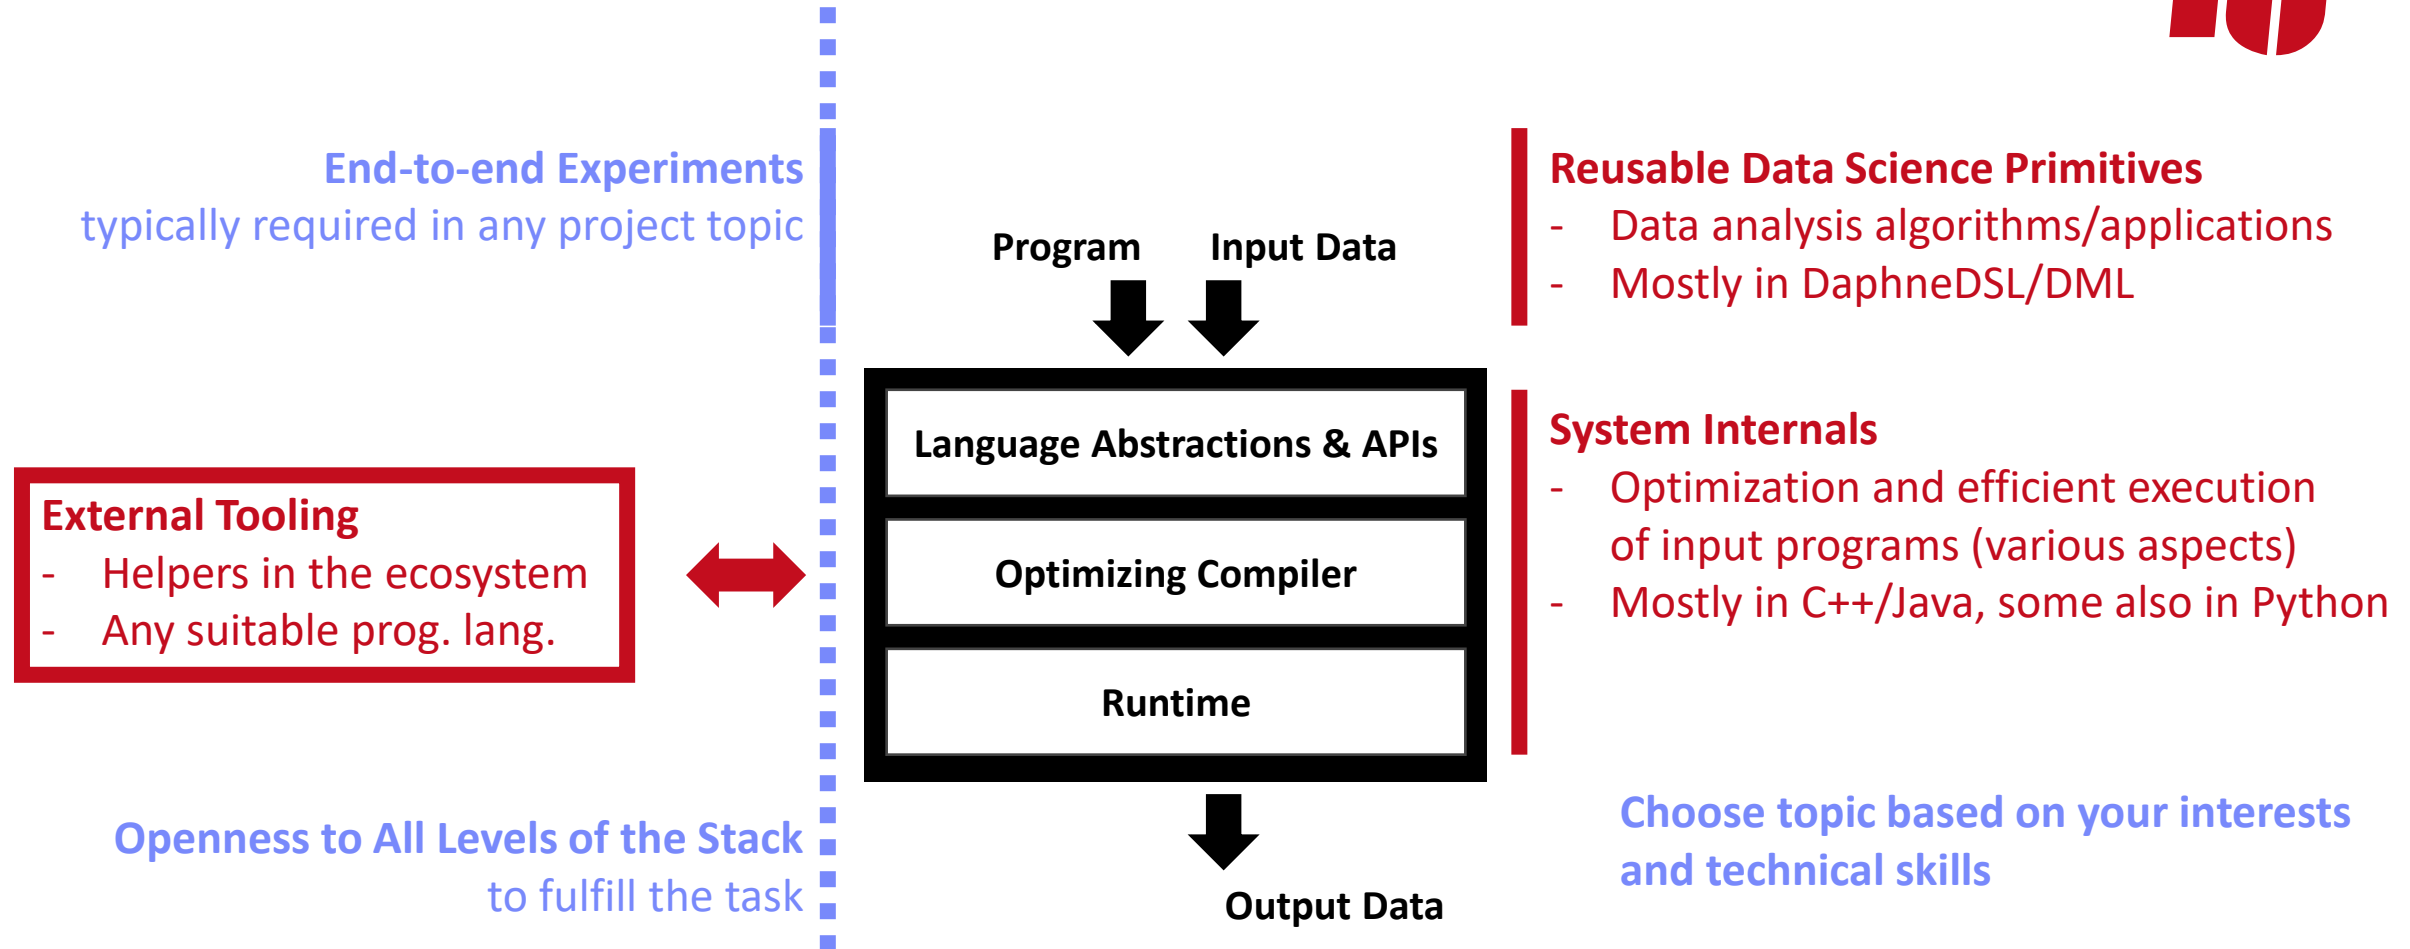
**Input Data**  
open file formats  
data exchange with other systems/libs



**Output Data**



# Kinds of LDE Project Topics



- **Individual/Team Project Work**

- Teams of up to 3 students **strongly encouraged**
- Unique topic for each individual/team

- **Ambitious Projects**

- 9 ECTS (~270 h of work)
- **≈6.75 weeks of full-time work**

- **Potential for Impact**

- Real open issues in existing systems
- If successful: meaningful contributions that will be used by others

- **Remarks on Topic Descriptions**

- Many open topics in DAPHNE and SystemDS
- Initial topic descriptions of varying level of detail
- During topic selection: Approach project mentor directly if interested in more details
- After topic assignment: More detailed descriptions where necessary
- We're open to alternative topic proposals

# How to Approach the Project


# Getting Started: Setting Up Your Development Environment




## Goals

- Build the system from source
- Successfully run the test suite

## Navigate to the GitHub Repos




<https://github.com/apache/systemds>




<https://github.com/daphne-eu/daphne>

## Clone, Build, Test According to the Documentation




<https://daphne-eu.github.io/daphne/GettingStarted/#quickstart-for-developers>



<https://apache.github.io/systemds/site/install.html>

## Know Where to Find the Documentation



<https://daphne-eu.github.io/daphne/>



<https://apache.github.io/systemds/>

# Getting Started: Preparing Your First Contribution

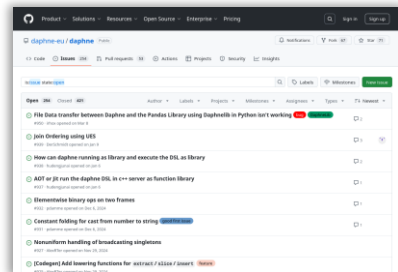


## ■ Goals

- Ability to modify the source code and run/test it
- Initial overview of relevant part of the code base

## ■ Set Up Your Editor/IDE etc.

## ■ Issue Tracking

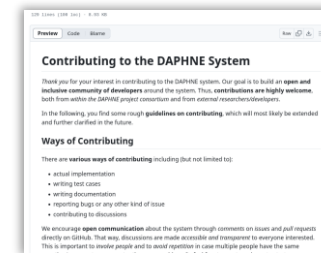


<https://github.com/daphne-eu/daphne/issues>

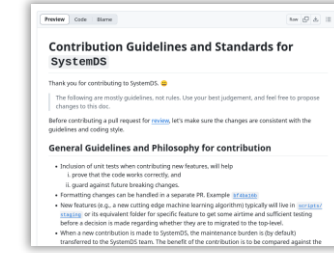


<https://issues.apache.org/jira/secure/Dashboard.jspa?selectPageId=12335852>

## ■ Read the Contribution Guidelines



<https://github.com/daphne-eu/daphne/blob/main/CONTRIBUTING.md>



<https://github.com/apache/systemds/blob/main/CONTRIBUTING.md>

## ■ Make Your First Modifications to the Code

- DAPHNE: “good first issues”
- SystemDS: Write initial test cases for your task

# Initial Prototype: Design/Implementation/Tests ( $\approx 2/3$ of the semester)



## ■ Goals

- 80% functionally complete prototype including good set of test cases
- Basis for further improvements driven by experiments and feedback

## ■ Mindset

- **Understand the topic**  
(task description, mentor, additional material)
- **Understand the code base**  
(overview and relevant parts)
- **Understand the employed libraries/frameworks**
- **Design and implement** step-by-step
- Not always “the” right solution: **explore alternatives**

## ■ Recommendations

- **Start as soon as possible**  
→ Don't underestimate the ramp-up effort
- **Actively approach your project mentor**  
→ Your mentor can give you valuable guidance

# Final Prototype: Creating a Convincing Contribution ( $\approx 1/3$ of the semester)



## ■ Goals

- High-quality code contribution whose value can easily be appreciated and understood

## ■ Recommendations

- **Submit tidy, well documented, extensively tested code**
- **View experiments as equally important as features,**
  - Start initial experiments as soon as possible
  - Focus primarily on investigating and improving your initial prototype after the intermediate presentation
  - **Experiments show the value of your contribution**
  - **You need time to incorporate your insights**
- Actively approach your project mentor**
  - Your mentor can give you valuable guidance



## ■ Goals

- Make your code easy to read/understand
- Others will have to maintain it after your contribution is merged

## ■ General Guidelines

- Clearly structure your code into **meaningful units** (classes, functions, etc.)
- Use **clear yet concise identifiers** (variable/function/class names)
- Stay **consistent with the existing code base** (e.g., use the same patterns) **OR refactor if necessary**
- Adhere to the code base's **coding style/formatting**

## ■ Keep Your Pull Request Tidy

- Stay focused: **Avoid changes unrelated to your task** (can be contributed as individual small pull requests)
- **Don't submit anything that's useless for others** (e.g., build artifacts, generated files (e.g., logs), IDE projects)
- **Exclude specifics of your local setup**: Avoid local paths, usernames, passwords, IP addresses etc.
- **Read your own pull request on GitHub** (changed files)

## ■ Goals

- Show functionally correct behavior
- Experiments don't make sense if prototype doesn't do what it should

## ■ General Hints

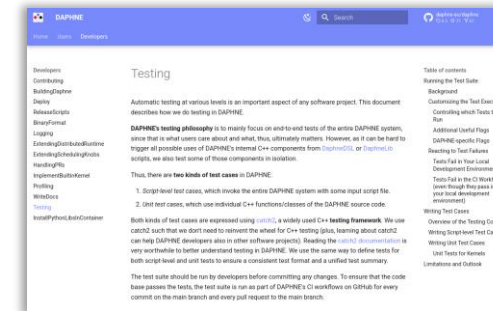
- **Unit tests** and **script-level tests**
- **Test cases that should work**
- **Test cases that should not work**  
(e.g., invalid DSL scripts, invalid input data, ...)
- Construct **simple and complex scenarios**
- Think of **corner cases**
- Small input data is often fine  
(but some bugs only triggered by large inputs)

## ■ Integrate with the Existing Test Suite

### ■ DAPHNE

→ directory: test/

→ see the documentation on writing test cases



<https://daphne-eu.github.io/daphne-development/Testing/>

### ■ SystemDS

→ directory: src/test/

# Final Prototype: Documentation



## ■ Goals

- Make your contribution understandable for **users** and **developers**

## ■ User Documentation

- **What's not documented doesn't exist**
- Add **high-level explanation** of features and concepts behind them
- Update documentation of **language abstractions** (e.g., new DSL built-in functions, new types, ...)
- Update documentation of **user APIs** (e.g., new command-line arguments)

## ■ Developer Documentation

- **Negative example: “The code is the documentation”**
- **High-level explanation** of your contribution; justify design decisions
- **API documentation** of classes, functions, members, etc.  
(integrate with the existing source code documentation of the system, e.g., doxygen or javadoc style)
- **Comments within function bodies** (e.g., high-level steps of an algorithm)

# Final Prototype: Experiments



## ■ Goals

- Understand your prototype to discover potential for improvement (e.g., performance bottlenecks)
- Demonstrate functional improvements (new features)
- Showcase non-functional improvements (performance compared to status quo and state-of-the-art baselines)

## ■ Design Experiments

- **Don't just conduct any random experiments**
- **Think about which questions you want/need to answer, design experiments accordingly**
- **Types of experiments:** Exploratory, micro benchmarks, benchmarks, end-to-end applications
- **Aspects of an experiment:** Data, workload, baselines, hardware & software stack, metrics

## ■ Conduct Experiments

- Automate as much as possible for repeatability (shell scripts etc.)

See 3<sup>rd</sup> seminar intro lecture on  
“Experiments, Reproducibility”

## ■ Visualize and Interpret the Results

- Automate the visualization based on raw experimental data
- Draw conclusions and react

# List of Project Topics (Proposals)

See list at [https://pdamme.github.io/teaching/2025-26\\_winter/lde/ProjectTopics.pdf](https://pdamme.github.io/teaching/2025-26_winter/lde/ProjectTopics.pdf)

# Summary and Q&A



- Course Organization, Outline, and Deliverables
- Projects in DAPHNE and Apache SystemDS
- How to Approach the Project
- List of Project Topics (Proposals)
- Remaining Questions?
- Reminder: Seminar Introductory Lecture Recommended for the Project
  - 03 Experiments, Reproducibility, and Giving Presentations [Oct 27, 14:00]
- See you during the consultation hours and intermediate presentations 😊